



НАЦИОНАЛЬНЫЙ ИССЛЕДОВАТЕЛЬСКИЙ
УНИВЕРСИТЕТ

Tokenisation and word segmentation

Francis M. Tyers

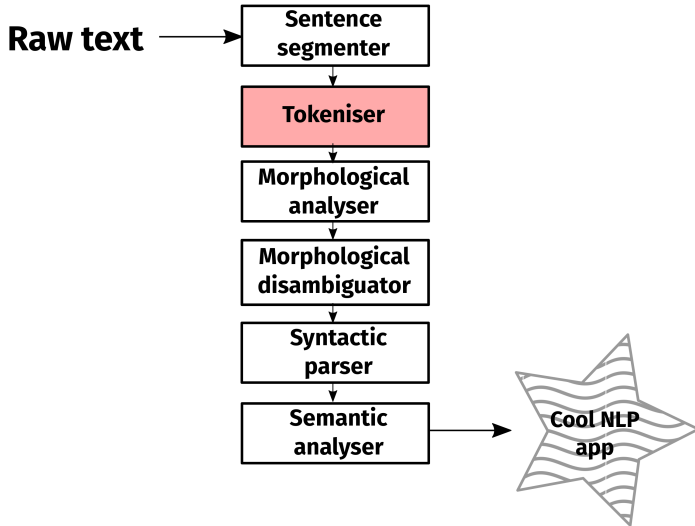
ftyers@hse.ru

<https://www.hse.ru/org/persons/209454856>

Национальный исследовательский университет
«Высшая школа экономики» (Москва)

1 октября 2018 г.

- Sentences



ARMA·VIRVMQVE·CANO·TROIAE·QVI·PRIMVS·AB·ORIS
ITALIAM·FATO·PROFVGVS·LAVINIAQVE·VENIT
LITORA·MVLTVM·ILLE·ET·TERRIS·IACTATVS·ET·ALTO
VI·SVPERVM·SAEVAE·MEMOREM·IVNONIS·OB·IRAM

- ...
- Vast majority of languages use some kind of whitespace-based word separator

Some questions:

- Multiword expressions
 - *только что* or *только-что* ?
- Named-entities
 - *Нижний Новгород* or *Нижний-Новгород* ?
- Numeral expressions
 - *150 000,0* or *150-000,0*

And how about abbreviations:

- *Гипотеза была выдвинута Каролом Борсуком в 1933 г.*
 - Is there one " " here or two?
- *В 1933 г. гипотеза была выдвинута Каролом Борсуком.*
 - And here?

Clitics:



The ideal tokenisation may depend on the task.¹

- Russian–Arabic MT:
 - Split off clitics²

¹ And also on the language (pair)!

² Zalmout and Habash (2017) “Optimizing Tokenization Choice for Machine Translation across Multiple Target Languages”

```
import sys, re

abbr = ['etc.', 'e.g.', 'i.e.']

def tokenise(line, abbr):
    line = re.sub(r'([\(\)\":?;!;])', r' \g<1> ', line)
    line = re.sub(r'([^\0-9])', r' \g<1> ', line)
    line = re.sub(r'([^\0-9])', r' \g<1> ', line)
    line = re.sub(r' +', ' ', line[:-1])
    output = []
    for token in line.split(' '):
        if token[-1] == '.' and token not in abbr:
            token = token[:-1] + ' .'

        output.append(token)

    return ' '.join(output)

line = sys.stdin.readline()

while line != '':
    print(tokenise(line.strip('\n'), abbr))
    line = sys.stdin.readline()
```

- Split off always-separating punctuation


```
import sys, re

abbr = ['etc.', 'e.g.', 'i.e.']

def tokenise(line, abbr):
    line = re.sub(r'([\(\)\":?;!;])', r' \g<1> ', line)
    line = re.sub(r'([^\0-9])', r' \g<1> ', line)
    line = re.sub(r'([^\0-9])', r' \g<1> ', line)
    line = re.sub(r' +', ' ', line[:-1])
    output = []
    for token in line.split(' '):
        if token[-1] == '.' and token not in abbr:
            token = token[:-1] + ' .'

        output.append(token)

    return ' '.join(output)

line = sys.stdin.readline()

while line != '':
    print(tokenise(line.strip('\n'), abbr))
    line = sys.stdin.readline()
```

- Split off commas not part of numeral expressions

```
import sys, re

abbr = ['etc.', 'e.g.', 'i.e.']

def tokenise(line, abbr):
    line = re.sub(r'([\(\)\":?;!;])', r' \g<1> ', line)
    line = re.sub(r'([^\0-9])', r' \g<1> ', line)
    line = re.sub(r'([^\0-9])', r' \g<1> ', line)
    line = re.sub(r' +', ' ', line[:-1])
    output = []
    for token in line.split(' '):
        if token[-1] == '.' and token not in abbr:
            token = token[:-1] + ' .'

        output.append(token)

    return ' '.join(output)

line = sys.stdin.readline()

while line != '':
    print(tokenise(line.strip('\n'), abbr))
    line = sys.stdin.readline()
```

- Collapse multiple spaces

```
import sys, re

abbr = ['etc.', 'e.g.', 'i.e.']

def tokenise(line, abbr):
    line = re.sub(r'([\(\)\":?;!;])', r' \g<1> ', line)
    line = re.sub(r'([^\0-9])', r' \g<1> ', line)
    line = re.sub(r'([^\0-9])', r' \g<1> ', line)
    line = re.sub(r' +', ' ', line[:-1])
    output = []
    for token in line.split(' '):
        if token[-1] == '.' and token not in abbr:
            token = token[:-1] + ' .'

        output.append(token)

    return ' '.join(output)

line = sys.stdin.readline()

while line != '':
    print(tokenise(line.strip('\n'), abbr))
    line = sys.stdin.readline()
```

- Split of full stops not part of abbreviations

- Let's meet at 17:45
- No one uses Yahoo! any more

- Some languages are written without spaces
- Tokenisation more difficult
- A number of algorithms available

- Rule-based algorithm
- Requires some kind of lexicon/dictionary
 - From a corpus
 - From a wordlist

function MAXMATCH(sentence, dictionary) **returns** word sequence W

if sentence is empty

return empty list

for $i \leftarrow \text{length}(\text{sentence})$ **downto** 1

firstword = first i chars of *sentence*

remainder = rest of *sentence*

if InDictionary(*firstword*, dictionary)

return list(*firstword*, MaxMatch(*remainder*, dictionary))

no word was found, so make a one-character word

firstword = first char of *sentence*

remainder = rest of *sentence*

return list(*firstword*, MaxMatch(*remainder*, dictionary))

- Start at beginning of string
- Iteratively look up the longest word in the dictionary
- If no word is found, output a single character

wecanonlyseeashortdistanceahead
we **canon**l y see **ash ort** distance ahead

– Alan Turing

- Works pretty well for some languages (e.g. Chinese)
- Not so great for others
- Why? Length of words

Word segmentation systems can be evaluated using Word Error Rate (WER):

- Edits (insertions, deletions, substitutions)

wecanonlyseeashortdistanceahead

Ref: we can only see a short distance ahead

Test: we can only see a short distance ahead

WER =



- Downside: requires an existing tokenised corpus





Watch out for characters other than space (U+0020):

- Non-breaking spaces (U+2060, U+FEFF, ...)
- Soft-hyphen (U+00AD)
- En quad (U+2000)
- En space (U+2002)

Plus 20 or so other characters.

Implement the `maxmatch` algorithm and test it on Japanese:

- Extract surface form dictionary from the training corpus
- Run the algorithm with the dictionary on the test corpus
- Write script to calculate WER for the segmentation.
 - Feel free to use a library for this

https://github.com/UniversalDependencies/UD_Japanese-GSD